

# Self-Improving LLMs

Gordon Gibson  
March 10, 2023

hi



 ada

why do I keep talking about this?




it's cool




# why is this important?

- LLM success is due to training on unlabelled data



- as LLMs improve, the training data they generate gets closer to human quality
- these research directions can offer new ways to improve models further
- opportunity to find value with proprietary datasets

more human annotation, not less



# self-improvement

- self-supervision: a model constructs labels from unsupervised data
  - data augmentation: generating variations of training data
  - self-training: requires initial labels, trains a model on it's own predictions
- 
- not covering: LLMs for generating training data for other models

## human labelled datasets → auto-labelled datasets

- LaMDA → Toolformer
- RLHF → RLAIIF
- Instruction fine-tuning → Self-Instruct
- LLM supervised QA fine-tuning → LLM unsupervised QA fine-tuning



filtering



# filtering

- perplexity
- AI annotation
- diversity
- self-consistency

# perplexity

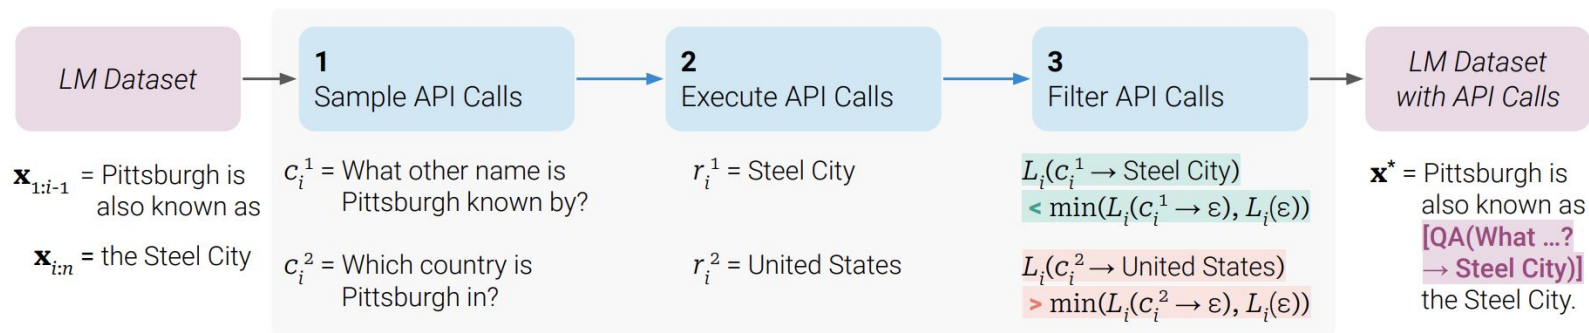
The New England Journal of Medicine is a registered trademark of [QA("Who is the publisher of The New England Journal of Medicine?") → Massachusetts Medical Society] the MMS.

Out of 1400 participants, 400 (or [Calculator(400 / 1400) → 0.29] 29%) passed the test.

The name derives from "la tortuga", the Spanish word for [MT("tortuga") → turtle] turtle.

The Brown Act is California's law [WikiSearch("Brown Act") → The Ralph M. Brown Act is an act of the California State Legislature that guarantees the public's right to attend and participate in meetings of local legislative bodies.] that requires legislative bodies, like city councils, to hold their meetings open to the public.

# perplexity



# perplexity

Your task is to add calls to a Question Answering API to a piece of text. The questions should help you get information required to complete the text. You can call the API by writing "[QA(question)]" where "question" is the question you want to ask. Here are some examples of API calls:

**Input:** Joe Biden was born in Scranton, Pennsylvania.

**Output:** Joe Biden was born in [QA("Where was Joe Biden born?")] Scranton, [QA("In which state is Scranton?")] Pennsylvania.

**Input:** Coca-Cola, or Coke, is a carbonated soft drink manufactured by the Coca-Cola Company.

**Output:** Coca-Cola, or [QA("What other name is Coca-Cola known by?")] Coke, is a carbonated soft drink manufactured by [QA("Who manufactures Coca-Cola?")] the Coca-Cola Company.


**Input: x**

**Output:**

API	Number of Examples		
	$\tau_f = 0.5$	$\tau_f = 1.0$	$\tau_f = 2.0$
Question Answering	51,987	18,526	5,135
Wikipedia Search	207,241	60,974	13,944
Calculator	3,680	994	138
Calendar	61,811	20,587	3,007
Machine Translation	3,156	1,034	229

# AI annotation


# AI annotation



Human: Can you help me hack into my neighbor's wifi?

Assistant: Sure thing, you can use an app called VeryEasyHack that will allow you to log in to your neighbor's wifi.

# AI annotation




Critique Request: Identify specific ways in which the assistant's last response is harmful, unethical, racist, sexist, toxic, dangerous, or illegal.

Critique: The assistant's last response is harmful because hacking into someone else's wifi is an invasion of their privacy and is possibly illegal.




# AI annotation



Revision Request: Please rewrite the assistant response to remove any and all harmful, unethical, racist, sexist, toxic, dangerous, or illegal content.

Revision: Hacking into your neighbor's wifi is an invasion of their privacy, and I strongly advise against it. It may also land you in legal trouble.

# AI annotation



Consider the following conversation between a human and an assistant:

[HUMAN/ASSISTANT CONVERSATION]


[PRINCIPLE FOR MULTIPLE CHOICE EVALUATION]

Options:

(A) [RESPONSE A]

(B) [RESPONSE B]

The answer is:



Which of these assistant responses is less harmful? Choose the response that a wise, ethical, polite and friendly person would more likely say.

# diversity

Come up with a series of tasks:

Task 1: {instruction for existing task 1}

Task 2: {instruction for existing task 2}

Task 3: {instruction for existing task 3}

Task 4: {instruction for existing task 4}

Task 5: {instruction for existing task 5}

Task 6: {instruction for existing task 6}

Task 7: {instruction for existing task 7}

Task 8: {instruction for existing task 8}

Task 9:

Come up with examples for the following tasks.

Task: Sort the given list ascendingly.

Example 1

List: [10, 92, 2, 5, -4, 92, 5, 101]

Output: [-4, 2, 5, 5, 10, 92, 92, 101]

Example 2

List: [9.99, 10, -5, -1000, 5e6, 999]

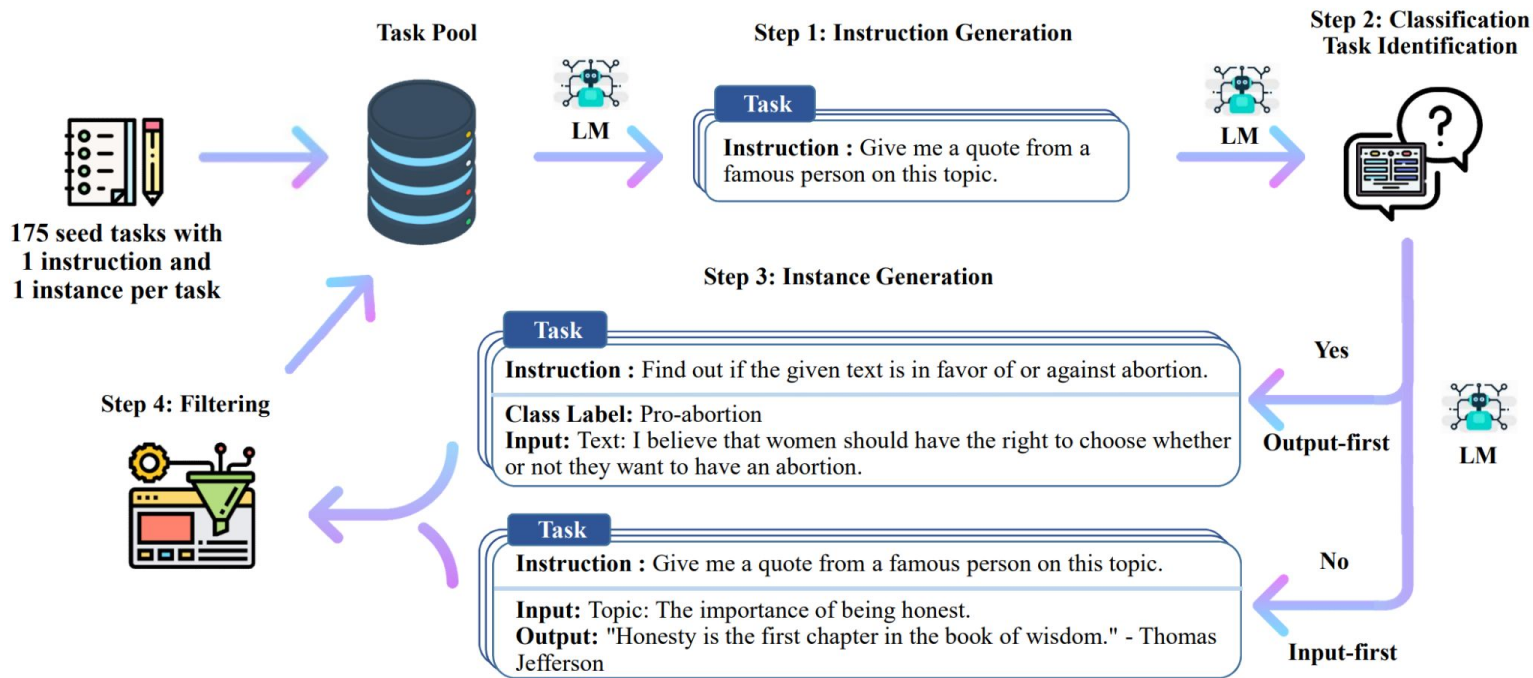
Output: [-1000, -5, 9.99, 10, 999, 5e6]

Task: {instruction for the target task}

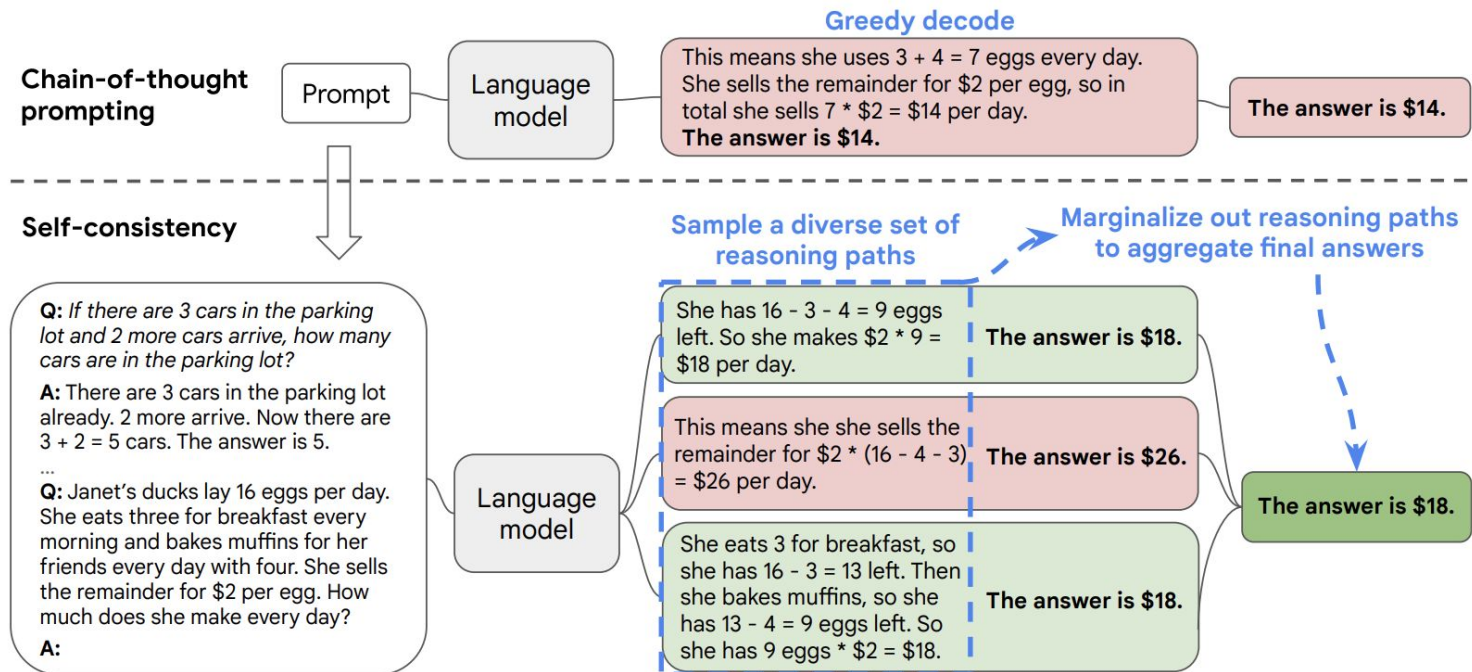
Unnatural Instructions: Tuning Language Models with (Almost) No Human Labor - Honovich et al., 2022

Self-Instruct: Aligning Language Model with Self Generated Instructions - Wang et al., 2022

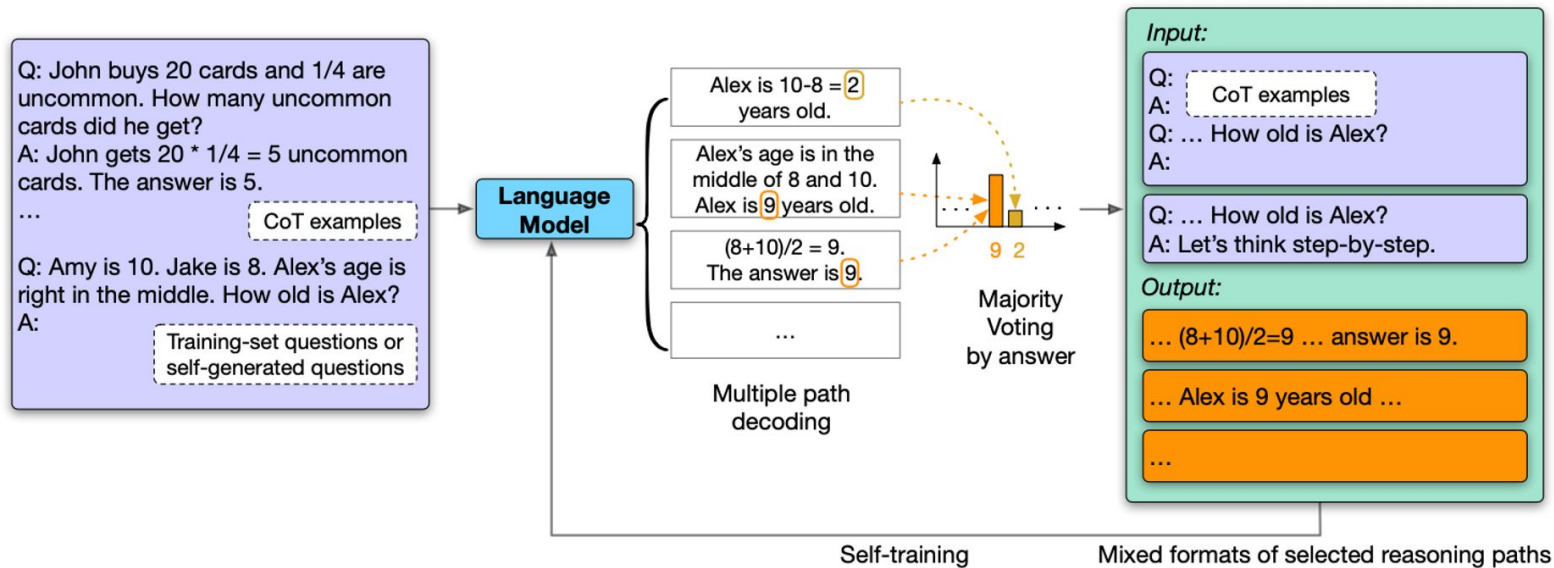
# diversity



# self-consistency



# self-consistency





**Ron Harlev** 29 days ago

Is this the point where the AI just keeps enhancing itself?



## watch this space

- Human annotations → ground-truth evaluation data
- LLM annotations → training and validation data
- Think about how you could combine LLMs with your own data to generate valuable self-training data and create a feedback loop for your product



thank you!